

How Are You Feeling? Analyzing Twitter User Mood with Naive Bayes Sentiment Analysis

Jesse Mu

December 12, 2014

1 Background

Sentiment analysis is a relatively new subset in the relatively new field of machine learning but has since found significant research interest and application in business analytics. Social networking site Twitter¹ has been identified as a key data source to assist in sentiment analysis study. Given its status as one of the biggest social networking sites in the world and its focus on trending topics and discussion using hashtags, Twitter users' tweets represent a massive, real-time snapshot of much of the world's opinions on countless topics—one of the most valuable and accessible resources of its kind.

Sentiment analysis is an example of a machine learning classification problem. As such, work in the field draws from standard classification models in the literature (e.g. Support Vector Machines, Maximum Entropy Models, Naive Bayes Classifiers) as applied to text classification [8]. The fundamental challenge of sentiment analysis is building a computer program that is able to “learn” from given data to generalize and form conclusions about new, unseen examples. This can be successfully approached by the more standard supervised (training data labeled with sentiment) approach or an unsupervised, lexical (unlabeled training data) approach [10].

The classical application of sentiment analysis is opinion mining of movie reviews, but the recent and dramatic surge in popularity of Twitter has resulted in substantial research interest in the potential for usage of English tweets as a corpus. Among the reasons for the usage of Twitter for sentiment analysis provided by Pak and Paroubek [7] are the inherent subjectivity of social microblogging platforms, a large and incredibly varied userbase, and potential for an “arbitrarily large” amounts of data.

When compared to other opinion mining tasks, sentiment analysis with a Twitter corpus stands out because of the linguistic properties of its training “documents.” Tweets are incredibly short messages, with a 140-character maximum limit, and encourage a fast, fundamentally electronic method of communication. The casual, spontaneous, and compressed nature of the service presents interesting and unique challenges. Specifically, feature extraction has been the distinguishing area of research when dealing with tweets, as researches encounter microblogging-specific problems including the usage of emoticons, abbreviations, sarcasm, and the sheer brevity of the information given.

¹<https://www.twitter.com/>

Twitter sentiment analysis is not limited to the academic space. A critical aspect of many companies' success is the successful monitoring and analysis of subjective information on a product through social media, and sentiment analysis on microblogging platforms plays a pivotal role. Because of the uniqueness of its content, the issue of Twitter sentiment analysis has led several companies to identify the problem as unique from standard sentiment analysis programs, and some text analysis companies such as Datumbox employ a separate Twitter analysis classification algorithm in addition to the traditional classifier. In addition to Datumbox's API, several other commercial services exist specifically for the sentiment analysis of English tweets. Many of them are at least partially free to use, including Sentiment140², Streamcrab³, and Semantria⁴. A well-performing sentiment analysis algorithm is a valuable asset, which motivates several of these companies to keep their code closed-source, like Sentiment140 and Semantria.

2 Problem Formulation

2.1 Corpus

Because Twitter sentiment analysis is a relatively new problem, there do not exist many good sources of labeled data for supervised learning tasks. It is possible to train a classifier based on standard sentiment data, e.g. from well-known movie review datasets, but the unique terminology and writing style of tweets demands a native training corpus.

Previous papers have used smaller training sets such as the 2013 SemEval corpus⁵ (~9000 tweets) and the Sanders Analytics training corpus⁶ (~5500 tweets). For this project, training data was taken from the Stanford University-associated online tool Sentiment140 mentioned previously, which provides a much larger corpus of 1,600,000 labeled tweets.

The sheer volume of tweets made available by Sentiment140 is due to the automatic collection of tweets based on a simplifying assumption: any tweet with positive emoticons (:D, :), =), etc) is labeled as a positive tweet, and any tweet with negative emoticons (: (, D:, =() is labeled as negative [3]. Intuitively, this assumption seems reasonably valid; due to the short nature of tweets, those with a positive emoticon are quite unlikely to have another negative emoticon resulting in an overall negative sentiment. Of course, there are certainly cases where this assumption fails, such as the usage of sarcasm and quoting. Another disadvantage of this data set is that it makes no qualification based on language; tweets of all languages are accepted, which decreases accuracy of a solely English-based classifier. However, the sheer volume of this data set outweighs other data sets by orders of magnitude and smooths out most anomalies.

²<http://www.sentiment140.com/>

³<http://www.streamcrab.com/>

⁴<https://semantria.com/>

⁵<http://www.cs.york.ac.uk/semeval-2013/task2/>

⁶<http://www.sananalytics.com/lab/twitter-sentiment/>

2.2 Naive Bayes

Because the main area of focus in Twitter opinion mining is good, reliable text preprocessing and feature extraction, multiple classifiers were not considered. The best classification algorithm for sentiment analysis tasks is disputed, and Go et al. [3] reports that the optimal model is unclear, but one common, well-performing model suggested by Gamallo [2], Narayanan [6] and Pak [7] and used commercially by companies like Datumbox is the Naive Bayes classifier.

For each document d , Naive Bayes attempts to assign the class c to d where

$$c = \max_{c_i} P(c_i|d) \quad (1)$$

Bayes' Rule shows us that for each c_i , $P(c_i|d)$ is equivalent to

$$P(c_i|d) = \frac{P(d|c_i)P(c_i)}{P(d)}. \quad (2)$$

The Naive Bayes algorithm then makes several simplifying assumptions. Because for each document d , $P(d)$ is constant, it can simply be ignored without consequence. Furthermore, $P(d|c_i)$ is calculated by considering the features of document d , x , where $x_1x_2x_3 \dots x_i$ are individual features. The Naive Bayes model *naively* assumes independence of these features, giving the simplified model

$$P(c_j|d) = P(c_j) \prod P(x_i|c_j). \quad (3)$$

Depending on whether or not we have a prior belief about the distribution of classes $P(c_j)$, the Naive Bayes model can either take this prior belief into account or learn from the training data class occurrences. If all classes are equally likely, the term $P(c_j)$ becomes irrelevant.

There are many advantages of the Naive Bayes classifier, although of course several disadvantages. Perhaps the benefit most tangible to this project is the speed and simplicity of the Naive Bayes algorithm, which requires considerably less time and memory consumption than similar classification algorithms such as SVM and ensemble learning methods [4]. Considering feature selection is paramount, a fast algorithm allows for rapid experimentation and prototyping, which makes detailed statistical analysis substantially easier. Despite its arguably oversimplifying assumptions, Naive Bayes is uniquely good at dealing with discrete text classification features, since more advanced feature selection techniques such as n-grams can compensate. Again, although disputed, Naive Bayes has been shown to outperform other algorithms in this task [6, 7].

Concerning disadvantages, one main issue with the Naive Bayes classifier is the lack of a reliable confidence metric when making decisions on new documents. Although the Naive Bayes classifier can output the probability it estimates for each class, the classifier tends to be overconfident [9]. This can pose problems when considering the web application in Section 5

There are two main variants of the Naive Bayes classifier used for discrete data, Bernoulli and Multinomial, which will be discussed further.

3 Technologies

The full source code for this project is open-source and available on GitHub at the following URL: <https://github.com/jayelm/twittersa>.

The code is written exclusively in Python, selected primarily due to software package availability. Text classification code uses the popular scikit-learn⁷ machine learning library which is built on the numpy⁸ and scipy⁹ scientific computing stack. The result is a heavily optimized machine learning library that retains the familiar and intuitive Python syntax while resorting to C bindings and optimizations when necessary for speed. It also incorporates many useful machine learning utilities to help automate feature extraction, pipelining, and evaluation.

The web application uses the web framework Flask¹⁰, and incorporates the scikit classifier module written for this project. Flask was chosen due to its minimalism, which suited the small project at hand; as a microframework, an entire web application can be written in only a few files, which is ideal for this project's small, single-use goal.

The cloud hosting service Heroku is used to host the web application. A development version is running at <http://twittersa.herokuapp.com/>.

There are also several utility scripts included in the project, written with Python and Shell, which assist in data sanitization and scraping.

4 Classifier Evaluation

The biggest portion of this project was the one specifically concerned with the natural language processing task at hand, and evaluation was integral to the development of my classification script and classifiers. Specifically, my goal was to consider several factors of a Naive Bayes classifier to identify a configuration most suitable for a production environment. To that end, I first and foremost attempted to maximize classification accuracy on test sets, while keeping in mind other important factors like portability (running time, memory constraints) and other metrics.

4.1 Preprocessing and Feature Vectorization

4.1.1 Preprocessing

Preprocessing refers to the regularization and normalization of documents to reduce extraneous document information like punctuation, stopwords, etc. in an attempt to eliminate the extraction of “noisy” features that can decrease accuracy and performance.

For this task, I created a `preprocess` function in `classifiers.py` that sanitized the document input. After tokenizing the input with NLTK's `word_tokenizer` method (more robust than scikit's), I implemented the following steps common in most text preprocessing implementations:

⁷<http://scikit-learn.org/stable/>

⁸<http://www.numpy.org/>

⁹<http://www.scipy.org/>

¹⁰<http://flask.pocoo.org/>

- Make all words lowercase.
- Decode strings into unicode literals.
- Stem each word with the NLTK Porter Stemmer.
- Remove punctuation, including Twitter-specific @ signs (“mentions”), # signs (“hash-tags”), and URL symbols.

There are also other issues unique to the language of Tweets that must be dealt with appropriately. Specifically, this includes the following implemented steps:

- Expanding acronyms (lol, wtf, lmao, etc) using a list of slang words from <http://www.noslang.com/>.
 - Since Noslang provides their dictionary freely on the website, but not in a computerized format, this required me to create the utility script `util/noslang_scraper.py`, which visits every page on the Noslang website, converting it into a Python dictionary, and serializing it.
- Removing the occurrences of two-or-more repeated characters to just two characters
 - E.g. “coool” → “cool”, “crazyyy” → “crazyy”
 - Duplicate characters cannot completely be removed since it’s too difficult to know when letters are duplicated due to the correct spelling of a word (e.g. “cool”) or emphasis (“crazyy”). Still, partial reduction is better than no reduction, and there may be value in differentiating between normal and emphasized uses of words.
- Emoticons are not considered in this training set. They are automatically removed by the Sentiment140 training corpus since an emoticon is present in every single training example such that they would carry no real information.
- Emoji (Japanese characters, of which many contain emoticons and other subjective images) are considered implicitly, when possible, since a subset of Emoji are represented by Unicode characters. Since I deal with unicode literals in my preprocessing step, these characters do have sentiments and counts associated with them.

After the text has been preprocessed, each document must be converted into a numerical vector to be processed efficiently by a classification algorithm.

4.1.2 The Bag-Of-Words Model

The most commonly used feature vectorization technique when dealing with text data is the bag-of-words model, which represents each document in a corpus as the unordered bag of its words.

Specifically, for a training corpus that contains n globally unique words, a unique numerical index is assigned to each word, 0 through $n - 1$. Once this global tally of words is created, either with a dictionary or by usage of a hash function, each document d is represented by

a feature vector $x \in \mathbb{R}^n$, where each element x_n stores the occurrence count of the n th word in the given document.

This model works very well when used in tandem with a Naive Bayes classifier, since such a classifier assumes independence of features and the bag-of-words model similarly doesn't keep track of word ordering.

Feature vectors created using the bag-of-words method can either be discrete counts or binary presence attributes, and results in different types of Naive Bayes classifiers that are discussed in Section 4.2.

Scikit-learn provides the `CountVectorizer` class for keeping track of the features and vectorizing a set of data. There are many feature extraction parameters that can be fine-tuned to produce better performance.

4.1.3 Laplace Smoothing

Laplace (+1) smoothing is implemented in the scikit Naive Bayes classifiers for dealing with unknown words. Lidstone (+0.5) smoothing was considered briefly as well, but appeared to have little to no effect on classifier metrics.

4.2 Initial Evaluation: Bernoulli versus Multinomial Naive Bayes

There are two main types of Naive Bayes classifiers that handle discrete data, and they make different assumptions about the probability distribution underlying the data. The first is the Multinomial Naive Bayes algorithm, which makes classification decisions assuming the features follow a multinomial (polynomial) distribution, while the second is the Bernoulli Naive Bayes, which makes classification decisions based on the Bernoulli (binary) probability distribution. The main difference between the two is that the Multinomial Naive Bayes considers the count of occurrences of a given word, rather than just its presence. On the other hand, the Bernoulli Naive Bayes model works only with binary features—it considers only the presence of given words. Scikit-learn provides two classes, `MultinomialNB` and `BernoulliNB` that implement these classifiers. Non-binary data given to a `BernoulliNB` instance is binarized and treated like a binary distribution.

One important decision to make before continuing with analysis is whether to use the Bernoulli or the Multinomial distribution. The easiest way to make this decision is to simply compare the performance of both on some naive, simple implementation of the classifiers.

For this initial evaluation, I used a standard bag-of-words approach with my preprocessing function. Training set size was broken into randomly sampled sets from 100 to 25000 tweets. A standard Bernoulli or Multinomial Naive Bayes classifier was run on a 70/30 train/test split, and was randomly sampled and tested 5 times to reduce the variance with test set selection. Figure 1 shows both the accuracy and F -score performance metrics for Bernoulli and Multinomial Naive Bayes classifiers based on varying training sets.

As evidenced from Figure 1, the difference between a Bernoulli and Multinomial distribution is very little, with the Bernoulli accuracy and F -score just slightly outperforming the multinomial metrics with data set of sizes 15,000+, with peak accuracy achieved by the Binomial Naive Bayes classifier at 74.21%. This is a good accuracy achieved with considering very standard implementation details.




Figure 1: Bernoulli vs Multinomial NB plotted against data set size

The small difference between the Bernoulli and Multinomial distributions is attributed to the short nature of Tweets. Multinomial classifiers take counts of words into account, with multiple words carrying higher weight than singular words. However, since Tweets are very short, multiple word occurrences do not occur often enough for sufficient information to be gained from considering multiple counts.

One other interesting thing to note is the effect of data size on classification accuracy. Although accuracy rises rapidly when dealing with smaller data sets (from $\sim 55\%$ accuracy to $\sim 75\%$), the classifier experiences diminishing returns when continuing to increase the data set size. From initial evaluation, performance on the classifier seemed to flatten out enough that a 25,000 tweet corpus was used as a good balance between accuracy and speed for the standard evaluation size for the rest of the experiments conducted here. However, as we will see later, greatly increasing data (by orders of magnitude) does continue to improve classifier performance.

Since a Bernoulli Naive Bayes distribution is mildly more efficient, we will continue using it alone, especially considering that the Bernoulli Naive Bayes classifier is generally recommended over the Multinomial for short documents [5].

4.3 Improving Classification Accuracy

The task at hand then becomes: how can we improve classification accuracy of the model developed? With the selection of a Bernoulli Naive Bayes classifier, the bulk of the potential for improvement lies in two main factors:

- Improving feature extraction and selection
- Obtaining more data

4.3.1 N-grams

The most basic NB classifier tested in the initial evaluation simply considers all globally unique n words as features. One way to potentially improve classification accuracy is by

including strings of words as features. For example, “The movie was not great” can be broken down into bigrams “The movie”, “movie was”, “was not”, and “not great”. Trigrams are similarly constructed

The theoretical underpinning behind N-grams is that longer, combined sections of words can capture more sophisticated language constructs than unigrams can. Most notably, the inclusion of bigrams as features is one solution to the negation problem, where the classifier can then consider directly negated features, i.e. “not good”. Of course, “not” and “good” are still considered as unigram features, but the idea is that the presence of the bigram would outweigh these two features. Furthermore, there is value in recognizing the difference between “somewhat good” and “extremely good”, two adjectives that would be missed from unigram feature selection but included and differentiated by bigrams.

These word-constructs are especially useful to Naive Bayes and bag-of-words classifiers since they assume independence between all features. By including ordered sections of words, N-grams add order considerations to the classifier, not in the classification algorithm itself but the feature selection. As displayed in Table 1, accuracy gains are relatively minimal, however, we see a jump of about 0.7% accuracy when including bigrams. Trigrams prove to overfit the data, and result in an accuracy reduction of 0.35%; this is probably due to the brevity of Tweets.

N-grams	Accuracy
1	0.7408
2	0.7473
3	0.7448

Table 1: Classifier accuracy for the cumulative addition of N-grams, 25000 data set, averaged over 5 samples

In scikit-learn, N-gram selection is implemented by setting the `ngram_range` variable during creation of the `CountVectorizer`. (1, 1) is the default, with cumulative bigrams and trigrams represented by (1, 2) and (1, 3) respectively.

4.3.2 Reducing features: select k -best

With the improved model considering bigrams in addition to unigrams, one of the ways commonly reported to increase accuracy is to actually *remove* features; in doing so, the classifier is able to eliminate “noisy” features. Selecting the k -best features in a data set is an attempt to statistically determine the most “informative” features, and use only those features.

“Noisy” features are those features that are probably irrelevant to the sentiment of a data, but information learned from a training corpus may still associate a given sentiment for a word; for example, due to the occurrences of the word “sandwich” in the training dataset, it may inadvertently carry a negative sentiment, although intuition tells us that “sandwich” should be neither a positive or a negative word. The goal of select k -best feature selection is to remove words like these that may skew document predictions.

There are multiple ways to evaluate which features are the most informative, including selection via mutual information, chi-squared statistical testing, ANOVA filtering, and more.

Manning [5] recommends the chi-squared (χ^2) statistical test for feature selection on text classification, which is implemented in scikit as `sklearn.feature_selection.chi2`, with the corresponding select k -best implementation at `sklearn.feature_selection.SelectKBest`. The `SelectKBest` instance takes a vectorized 2D array of feature vectors, where it applies a given statistical test to determine the user-supplied k -best features, trimming the array down into an $m \times k$ matrix.

For the 25,000 tweet data set, there are around 140,000 unigram + bigram features. Figure 2 shows the effect of select k -best feature selection on classifier accuracy, averaged over 5 splits, based on varying amounts of k . Surprisingly, select k -best feature selection does not appear to have a significant effect on classifier accuracy. Obviously, when the number of features selected is very low (e.g. 10), there is not nearly enough information to make reliable decisions on the data, and the accuracy is very low. However, anywhere in the area between 5,000 and 140,000 features, select k -best feature selection has little-to-no effect on classifier accuracy.

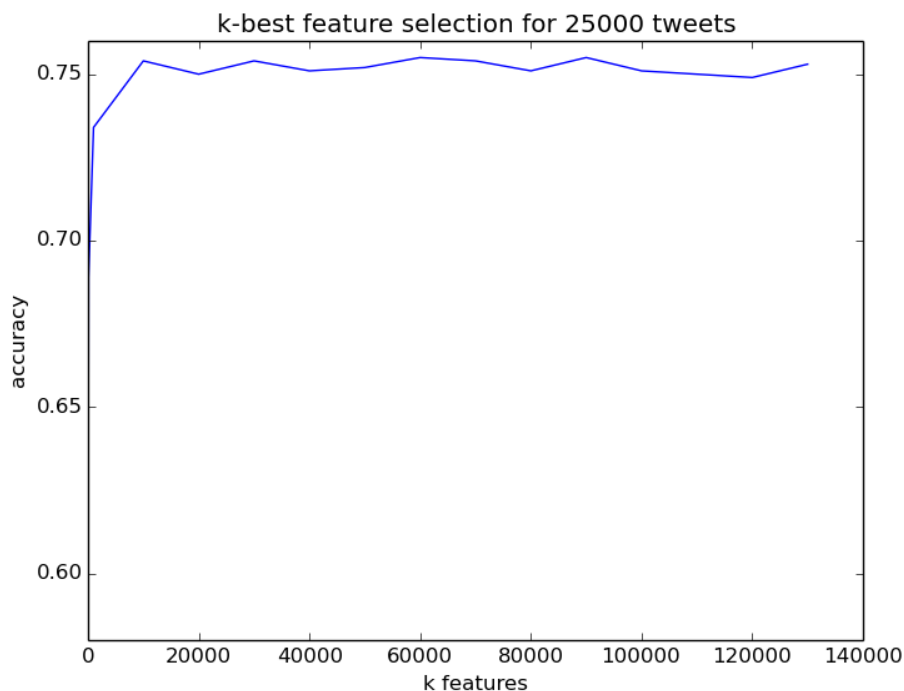


Figure 2: Select k -best feature selection performance on 25000 data set (bigrams + unigrams), averaged over 5 splits

This is probably due to the fact that for bigger data sets and an equal distribution of positive and negative cases, irrelevant words tend to appear equally among both cases (i.e. the word “the” and “sandwich” should both end up appearing in around 50% positive and negative cases). The Naive Bayes classifier is mathematically insensitive to irrelevant features [5], so such features would not have a big effect on classifier accuracy. Only on somewhat smaller data set sizes would select k -best feature selection make a bigger difference in classifier accuracy.

4.3.3 Other Feature Selection Techniques

Both “term inverse, inverse document frequency” (TF-IDF) and variance threshold removal feature selection techniques were considered, but neither appeared to have any tangible effect on classifier accuracy. TF-IDF is a vector-transformation method that increases the importance of words depending on how often they appear in a given document (TF) but downplays the importance of words that appear several times in all documents (IDF). Because of the brevity of the words in the Tweet corpus, TF-IDF does not have much effect on classifier accuracy.

Variance threshold feature removal removes features whose variance does not pass a certain threshold. Multiple thresholds were tested (0, 0.5, 1), but none seemed to have a large effect.

4.3.4 More Data - Diminishing Returns

One of the use cases in which a Naive Bayes classifier is recommended is when large amounts of data is available. A Naive Bayes classifier is crude but fast, so it can take advantage of large amounts of data while still maintaining a reasonable running time over comparable algorithms [1]. As it turns out, the single most effective way to increase test set accuracy on this project is to greatly increase the amount of data. Although 25,000 was used as the benchmark, Figure 3 shows the accuracy of a classifier (with bigrams and 0.0 variance threshold feature removal) on data sets at 25,000, 100,000, 250,000 and 500,000 tweets, averaged over 5 samples. There is marked increase in classifier accuracy dependent on increasing size, but at 500,000 tweets, memory and training time becomes an issue—the 500,000 tweet training set takes nearly 10 minutes to run and evaluate. A peak is achieved for the 500,000 tweet data set at 79.12%.

5 Building the Web App

5.1 Backend

As previously mentioned, the web application was written in Flask, and simply includes an index page, a user endpoint page, and an error page. Most of the code is included in `twittersa.py`. When submitting a request to the `/search?q=<USER-ID>` endpoint, the user function takes the supplied Twitter username and attempts to get historical tweets associated with that user by calling the `statuses/user_timeline` Twitter API endpoint. If no tweets are found, either because the user does not exist or the timeline is private, an error page is shown.

How many tweets the application gets is dependent on the `USER_API_CALLS` global variable set in `twittersa.py`. The API enforces a limit of 200 tweets per user timeline call, and keeps a maximum of up to 3200 tweets of historical data. Thus, this variable can be anywhere from 1 to 16, and will attempt to get as many tweets as possible given the API call limit.

The `status/user_timeline` API endpoint is rate-limited to 300 requests per minute for application-level authentication, so increasing `USER_API_CALLS` results in quicker exhaustion of the rate allowance. It is currently set at 2. The application currently uses a singular set of




Figure 3: Effect of large training sets on classifier accuracy

Twitter API keys from an application registered at dev.twitter.com, since I do not anticipate usage to be high enough such that the rate limit actually limits usage of the application.

After tweets are collected, they are sorted by date and put into histogram-like bins. Depending on the frequency of tweets, they are put into either weekly or monthly time intervals. Then, the production classifier specified from `sentiment/classifiers.py` predicts the sentiment for each tweet, returning an array of `TweetSentiment` objects. The average sentiment for each time interval is calculated, and transformed into a JSON-like object that will be used for display by the frontend. Both this object and a list of the original tweets are passed to the HTML view using the `jinja2` templating engine.

5.2 Frontend

`Chart.js`¹¹ is used as the visualization library for the tweet information. It was selected over comparable libraries for simplicity and ease of use, although it lacks the extensible functionality of libraries like `d3.js`¹². The Python object containing the `TweetSentiment` object is JSONified into a format understood by the library, which creates the line plot plotting average Tweet sentiment against each time interval. All tweets by the user are then displayed below the chart with dates and individual sentiment probabilities. Some additional JavaScript scripting is used to enable functionality, which is currently limited to clicking on

¹¹<http://www.chartjs.org/>

¹²<http://d3js.org/>

an individual node on the line chart and seeing the tweets specific to that date range filtered under the chart. An example screenshot of this library in action is located in Figure 4.

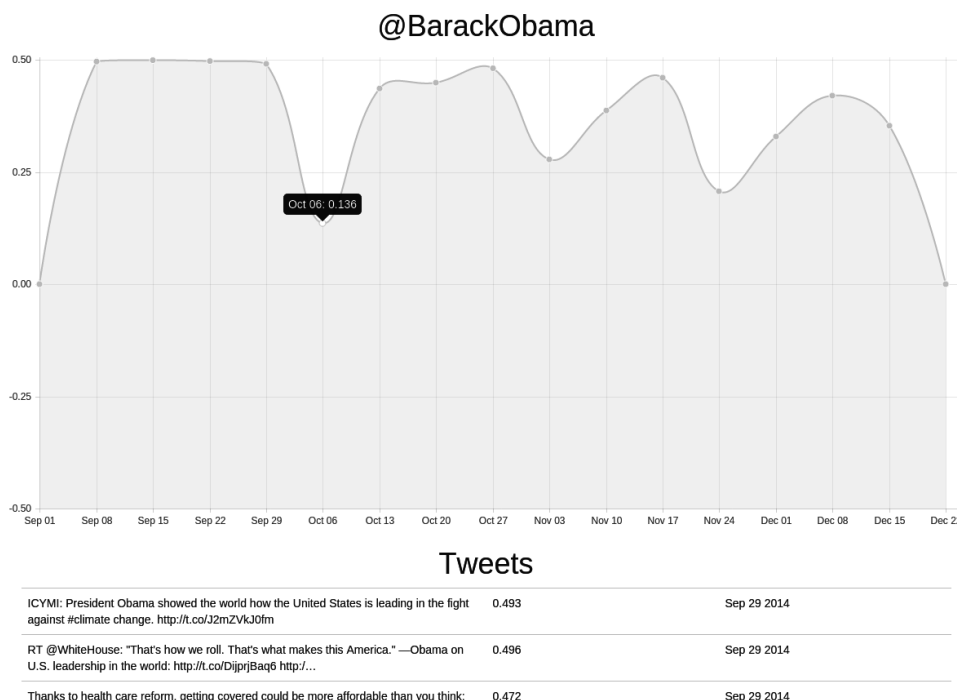


Figure 4: Example Web Application Screenshot for Twitter User @BarackObama

6 Conclusions

Using a standard Bernoulli Naive Bayes classification algorithm on the largest 500,000 tweet dataset (350,000/150,000 train/test split) with bigram feature selection, the peak accuracy achieved by this classification system is 79.12%. This is a surprisingly good result that shows the efficacy of simple, probabilistic machine learning techniques when applied to sentiment analysis.

Most notably, the task of sentiment analysis itself is not a 100% agreeable problem. Specifically, Wilson et al. [11], in an empirical analysis, identified that when participants were asked to classify sentences as “positive”, “neutral”, or “negative”, only 82% similarity was recorded, not very far above the accuracy level recorded by this classifier. When neutral tweets are removed, however, Wilson reported that accuracy increased to $\sim 90\%$. An 80% benchmark by a machine, however, is still quite good, and is comparable to the state-of-the-art work in statistical sentiment analysis; Pang et al. [8] report a maximum accuracy percentage with a Naive Bayes classifier of 81.0%.

There were of course several limitations to this study. Most of them were mentioned during the writeup as I consider the implications of various decisions used during the task. For example, the Sentiment140 dataset’s collection assumption could be inherently problematic (they are not hand-classified), but the size of the training set outweighs this possible

disadvantage. Also, training data was taken from this same corpus. As such, emoticons could not be considered, since they are removed from the original Sentiment140 training set.

Since a Naive Bayes classifier is somewhat overconfident, using its probability estimation as a “confidence” metric is somewhat problematic. This is somewhat mitigated by the fact that the Flask application averages all of the probabilities to gain a more global sentiment, but it is one of the inherent weaknesses in a Naive Bayes classifier.

Notably, many feature selection techniques that were considered to work well for text classification did not (especially the select k -best feature selection). This could be indicative of an inherent issue with the subject matter of the training corpus, and may need to be investigated further. Testing on different, hand-tagged training sets, for example, may yield better results.

Ideas for Further Work Only a binary classifier was built, which distinguishes between positive and negative examples. Since a NB classifier is somewhat overconfident on decisions, including a neutral class may actually improve accuracy. This could also improve the handling of decidedly neutral tweets (e.g. “I am eating a sandwich”). Identifying neutrality could be done by making a probability confidence threshold for polarity in the Naive Bayes classifier or other methods implemented by Gamallo [2] which first considers whether or not a tweet is Subjective or Objective before determining polarity.

With a bigger time investment, one interesting thing to do would be to look at precision and recall metrics along with simple accuracy. This could give a more sophisticated view of the performance of the classifier based on the different properties selected.

More feature selection and preprocessing techniques could be considered; for example, URLs can be normalized, instead of simply being stripped with punctuation, and mentions and hashtags could simply be completely removed instead of stripped of punctuation. A close analysis of accuracy would be necessary to determine whether these changes would be beneficial for the classifier. Narayanan [6] also implements a negation preprocessing step that prefixes negated words with `not_`. Although I implicitly try to amend this issue with N-gram feature selection, the negation approach is certainly more sophisticated and can probably result in better analysis.

It also may be possible to include some discrete features separate from the bag-of-words representation, including document length, number of “polarized” words according to a polarity lexicon, etc. These ideas are considered by Gamallo [2]. For tweets specifically, heuristic document features like document length will probably not play a large role, although a more content-based approach could (with appropriate feature scaling) play a role.

And again, with more time and memory, further increasing data (up to Sentiment140’s 1,600,000 tweet corpus) will most likely further increase accuracy. It may also be appropriate to examine other classification algorithms, such as SVM or Ensemble methods, and make a comparative analysis.

A Appendix: Python script help output

A.1 classifiers.py

```
usage: classifiers.py [-h] [-a] [-c CLASSIFIER] [-s] [-n NGRAM]
                    [--tf | --tfidf] [--show-best-features] [-k KBEST]
                    [-v [THRESHOLD]] [-p [PICKLE]] [-N NUM] [-r]
                    [corpus [corpus ...]]

positional arguments:
  corpus                corpus sizes (must be the size of a .pickle file in
                        lib/)

optional arguments:
  -h, --help            show this help message and exit
  -a, --all             grab all training .pickle files in lib
  -c CLASSIFIER, --classifier CLASSIFIER
                        specify classifier (bernoulli or multinomial)
  -s, --stopwords       filter out stopwords before processing
  -n NGRAM, --ngram NGRAM
                        use ngrams in addition to unigrams
  --tf                  use tf feature normalization
  --tfidf               use tf-idf feature normalization (incompatible with
                        --tf)
  -k KBEST, --k-best KBEST
                        select k best features with chi-squared statistical
                        test
  -v [THRESHOLD], --variance-threshold [THRESHOLD]
                        remove features with variance below threshold
  -p [PICKLE], --pickle [PICKLE]
                        save last classifier to given destination
  -N NUM, --num NUM     number of times to train each classifier
  -r, --repl            enter REPL for last classifier
```

A.2 twittersa.py

Note: environment variables must be set before running the script.

```
[2014-12-12 16:10:35,371] INFO in twittersa: Loading classifier...
[2014-12-12 16:10:57,365] INFO in twittersa: Done
usage: twittersa.py [-h] [-d]
```

```
optional arguments:
  -h, --help    show this help message and exit
  -d, --debug   run application in debug mode
```

References

- [1] Davis, E. Naive Bayes for Classifying Text. Lecture notes from *Introduction to Artificial Intelligence* (V22.0480.001, 2013), New York University. Retrieved from <http://www.cs.nyu.edu/faculty/davise/ai/>.

- [2] Gamallo, P., and Garcia, M. Citius: A Naive-Bayes Strategy for Sentiment Analysis on English Tweets. In *Proceedings of the 8th International Workshop on Semantic Evaluation* (Dublin, Ireland, 2014), 171–175.
- [3] Go, A., Bhayani, R., and Huang, L. Twitter Sentiment Classification using Distance Supervision. In *Processing* (2009), 1–6.
- [4] Huang, J., Lu, J., and Ling, C. X. Comparing Naive Bayes, Decision Trees, and SVM with AUC and Accuracy. In *The Third IEEE International Conference on Data Mining* (Melbourne, FL, 2003), 553–556.
- [5] Manning, C., Raghavan, P., and Schtze, H. /emphIntroduction to Information Retrieval. Cambridge University Press, Cambridge, 2008.
- [6] Narayanan, V., Arora, I., and Bhatia, A. Fast and accurate sentiment classification using an enhanced Naive Bayes model. In *Intelligent Data Engineering and Automated Learning Lecture Notes in Computer Science* (2013), 194–201.
- [7] Pak, A., and Paroubek, P. Twitter as a Corpus for Sentiment Analysis and Opinion Mining. In *LREC* (Valletta, Malta, 2010), European Language Resources Association, 1320–1326.
- [8] Pang, B., Lee, L., and Vaithyanathan, S. Thumbs up? Sentiment Classification using Machine Learning Techniques. in *Proceedings of EMNLP* (Philadelphia, PA, 2002), Association for Computational Linguistics, 79–86.
- [9] Rennie, J. Improving Multi-class Text Classification with Naive Bayes. Master’s Dissertation (Cambridge, MA, 2001), Massachusetts Institute of Technology.
- [10] Turney, P. Thumbs Up or Thumbs Down? Semantic Orientation Applied to Unsupervised Classification of Reviews. in *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics* (Stroudsburg, PA, 2002), Association for Computational Linguistics, 417–424.
- [11] Wilson, T., Wiebe, J., and Hoffman, P. Recognizing Contextual Polarity in Phrase-Level Sentiment Analysis. In *Proceedings of HLT/EMNLP* (2005), Association for Computational Linguistics, 347–354.